KEWEN WANG

Email: wangkewen001@gmail.com   Website: http://wangkewen.github.io
Address: Jersey City, NJ

EDUCATION

University of Connecticut   *2014 -2020*
Ph.D. in Computer Science, GPA: 4.0
Beihang University   *2010 -2013*
M.S. in Computer Science, GPA: 3.3
Beijing Information Science and Technology University   *2005 -2009*
B.S. in Computer Science, GPA: 3.5

TECHNICAL SKILLS

Computer Languages	Java, Python, Linux Shell, C, Go
Open Source	Apache Spark, Apache Hadoop, Apache Mesos, Aurora, MySQL, Postgres
Web Development	ReactJS, Node.js, JavaScript, Apache Tomcat

CODING COMPETITION

Google Code Jam 2017	Qualification Round Rank#1483/25k, Round 1C Rank#1664/3775
Google Kickstart 2018	Round B Rank #122/753

RESEARCH PROJECTS

Performance Prediction and Improvement for Apache Spark Jobs   Aug 2014 - May 2019
Research Assistant   *University of Connecticut*

· Developed a Spark analytics system in Java to parse JSON logs of Apache Spark event, and predict time, I/O overhead, memory consumption using analytical approaches.
· Developed a dynamical job predictor in Java to predict the execution time of multiple Spark jobs in Xen, and implemented a job scheduler in Java and Bash to reduce the total execution time.
· Implemented a Spark optimizer in Java to predict and mitigate potential task stragglers and skewed task distribution problems for Apache Spark platform to improve job performance.
· Designed and implemented a middleware to dynamically allocate computing resources for Apache Spark applications to improve resource utilization.

Optimizing Hadoop MapReduce   Nov 2011 - Dec 2012
Research Assistant   *Beihang University*

· Applied BTrace to trace MapReduce job functions, and monitor resource consumption using Ganglia.
· Implemented a MapReduce optimizer in Java through constructing Hadoop performance model for execution time prediction and designing heuristic search algorithm to find near optimal configurations for MapReduce jobs.

WORK EXPERIENCE

Software Engineer II   July 2019 - Current
Oscar Health. New York, NY

· Working in Compute Infrastructure team.
· Building platform services.

Research Intern   May 2018 - Aug 2018
HashiCorp. San Francisco, CA

· Developed a system performance predictor in Python for Consul cluster workload prediction using Machine Learning algorithms such as SVM, Random Forest, Gradient Boosting Tree.
· Implemented server buffer in Go to improve cluster stability and reduce response latency and failure.

Full Stack Developer   Nov 2011 - Jan 2012
Institute of Science and Technology at Beihang University

· Designed and implemented a project management system on Struts+Spring+Hibernate framework.
· Implemented information retrieval and display using JSP, JavaScript and Ajax, and loaded project archives into MySQL database
· Implemented Java Servlets and filters for service actions (such as information update, remove) and security validation (such as user access control), applied Apache Tomcat as the web server in Linux.

Software Engineer Intern   Mar 2010 - May 2010
NDtech Inc. Beijing, China

· Analyzed ANTLR (an open source parser generator) to learn C# parser and Script#.
· Applied Script# to write JavaScript using C#.

AWARDS

Predoctoral Fellowship   *2017*
Computer Science and Engineering department at University of Connecticut
Third Class Scholarship   *2011*
Beihang University
Academic Scholarship   *2008*
Beijing Information Science and Technology University
Municipal 2nd Prize of 21st National Middle School Students Physics Competition *2004*
City of Xianning, China

PUBLICATIONS

1. A Dynamic Resource Allocation Framework for Apache Spark Applications. Wang, Kewen, Mohammad Maifi Hasan Khan, and Nhan Nguyen. 2020 IEEE 44th Annual Computer Software and Applications Conference (COMPSAC).

2. A Model Driven Approach towards Improving the Performance of Apache Spark Applications. Wang, Kewen, Mohammad Maifi Hasan Khan, Nhan Nguyen, and Swapna Gokhale. 2019 IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS).

3. Modeling Interference for Apache Spark Jobs. Wang, Kewen, Mohammad Maifi Hasan Khan, Nhan Nguyen, and Swapna Gokhale. IEEE 9th International Conference on Cloud Computing (CLOUD), 2016.

4. Performance Prediction for Apache Spark Platform. Wang, Kewen, Mohammad Maifi Hasan Khan. IEEE 17th International Conference on High Performance and Communications (HPCC), 2015.

5. Predator - An experience guided configuration optimizer for Hadoop MapReduce. Wang, Kewen, Xuelian Lin, and Wenzhong Tang. IEEE 4th International Conference on Cloud Computing Technology and Science (CloudCom), 2012.